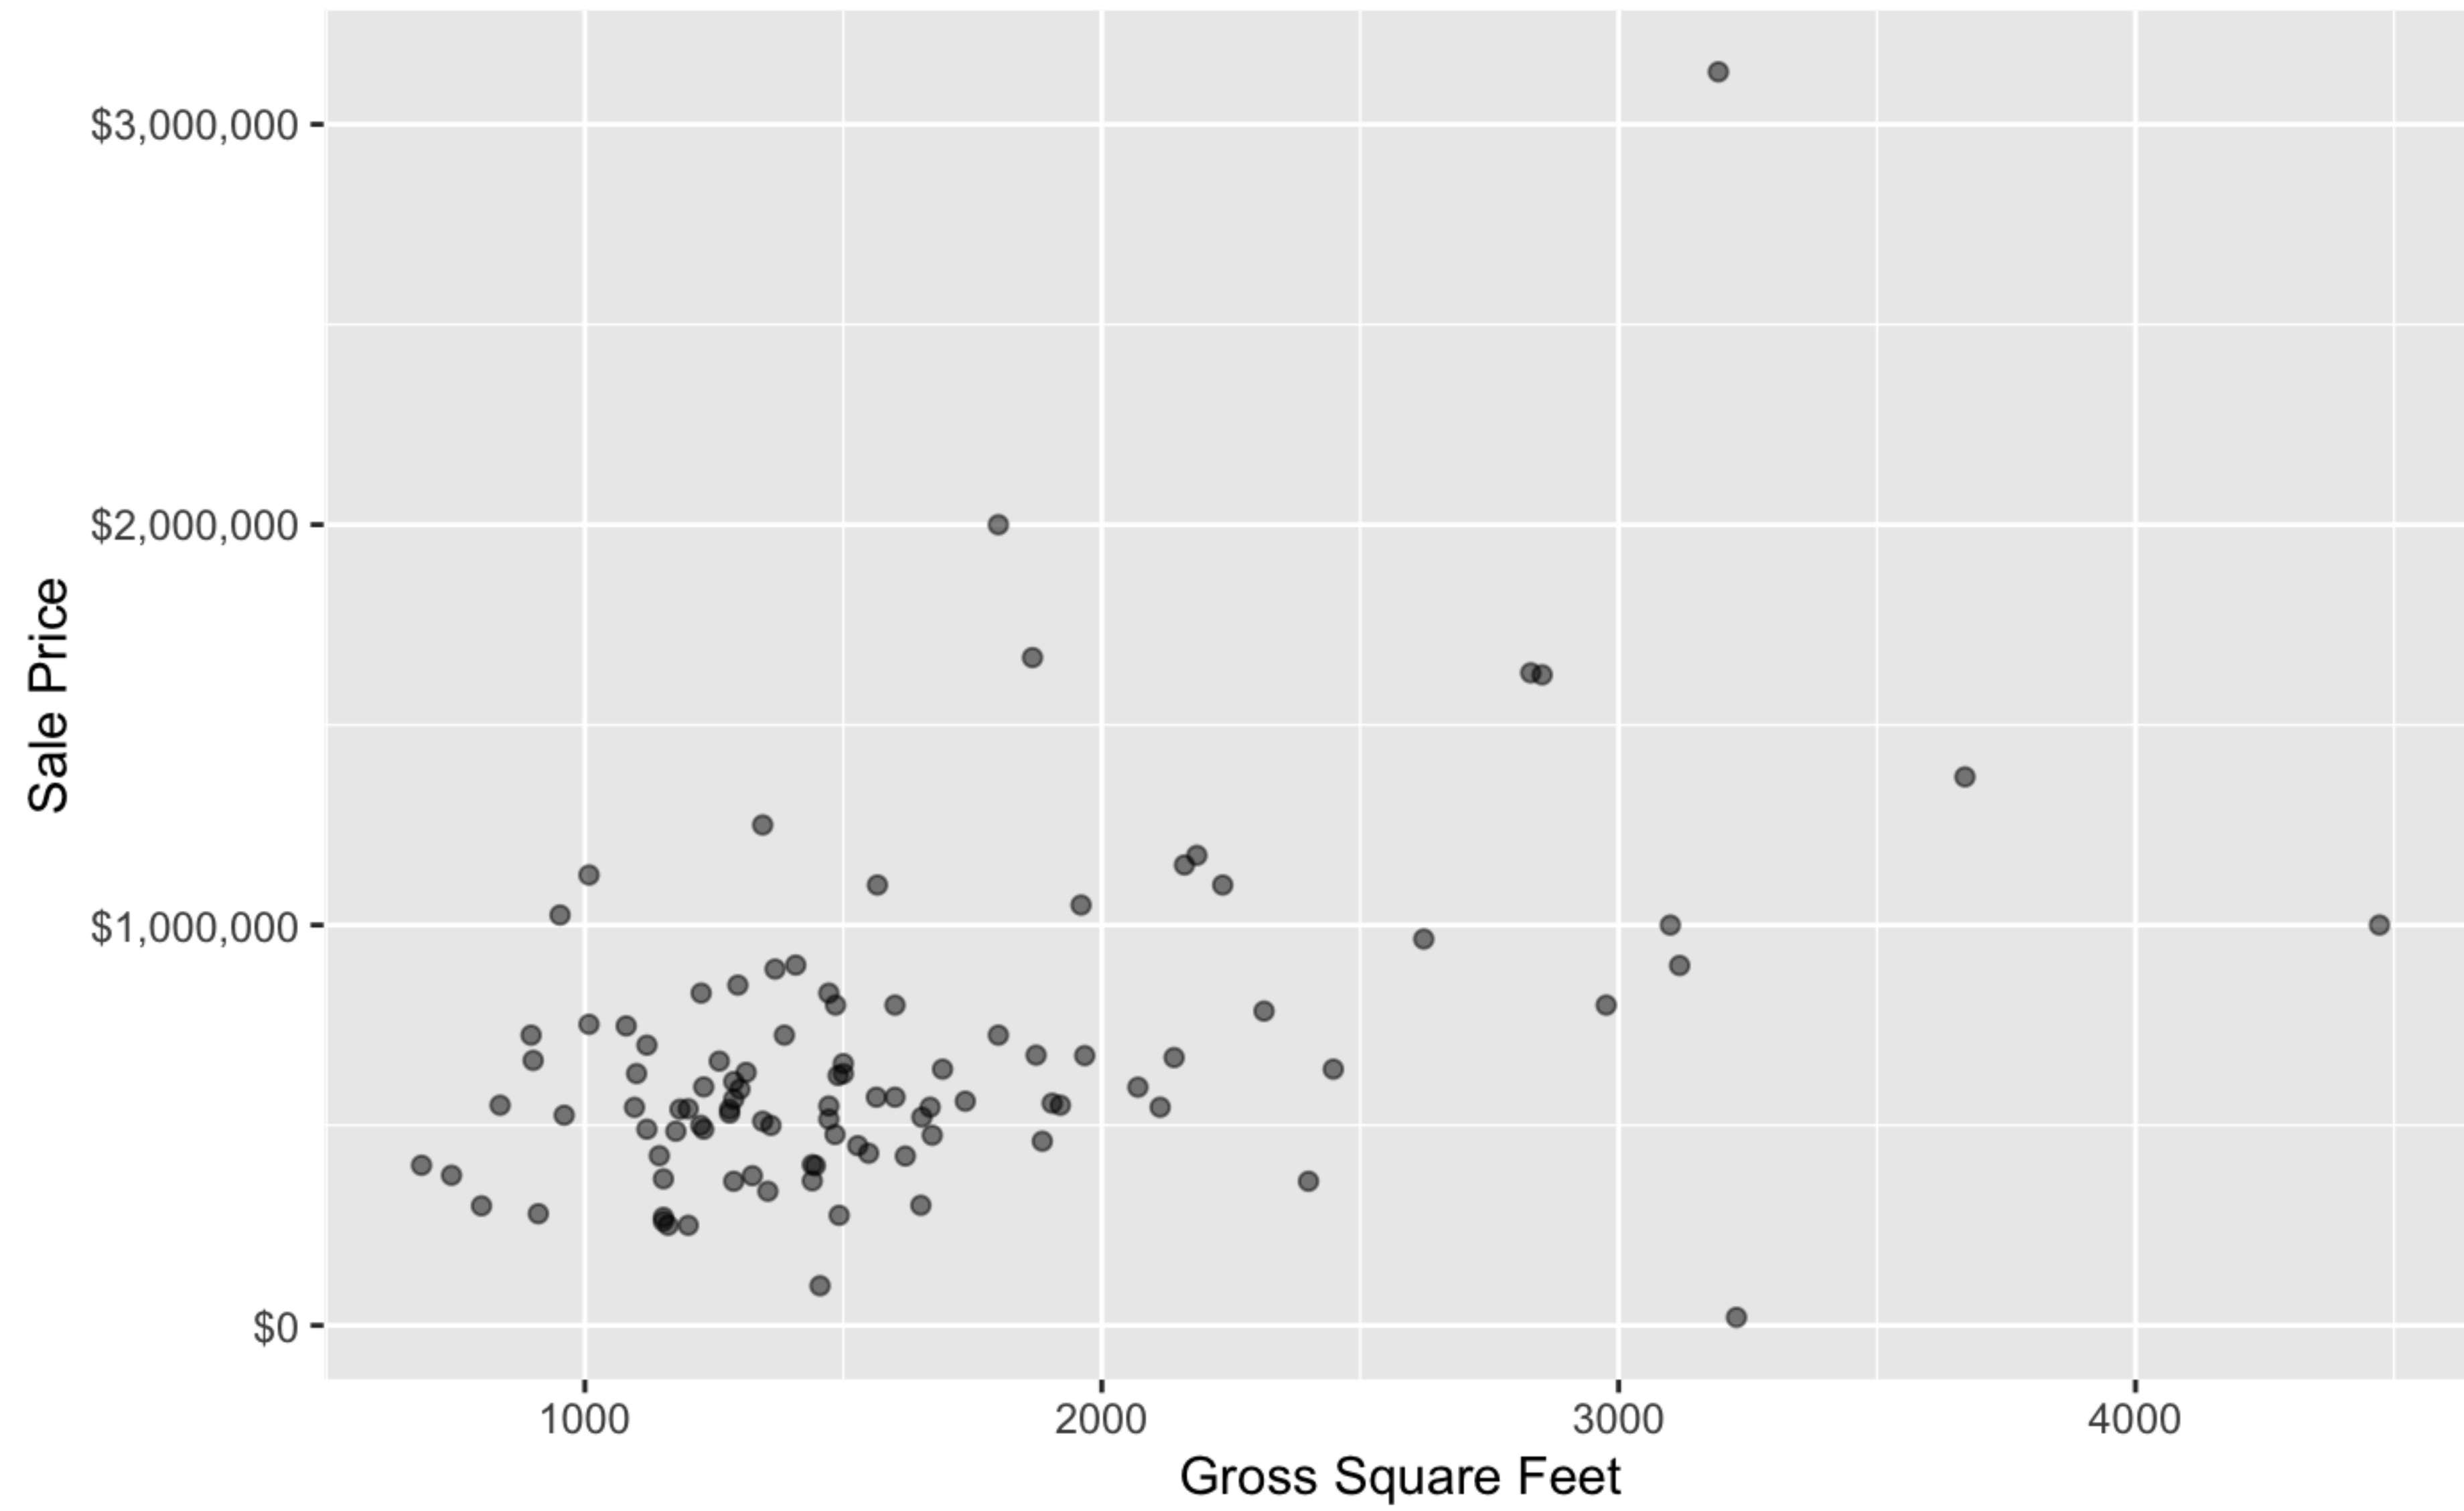



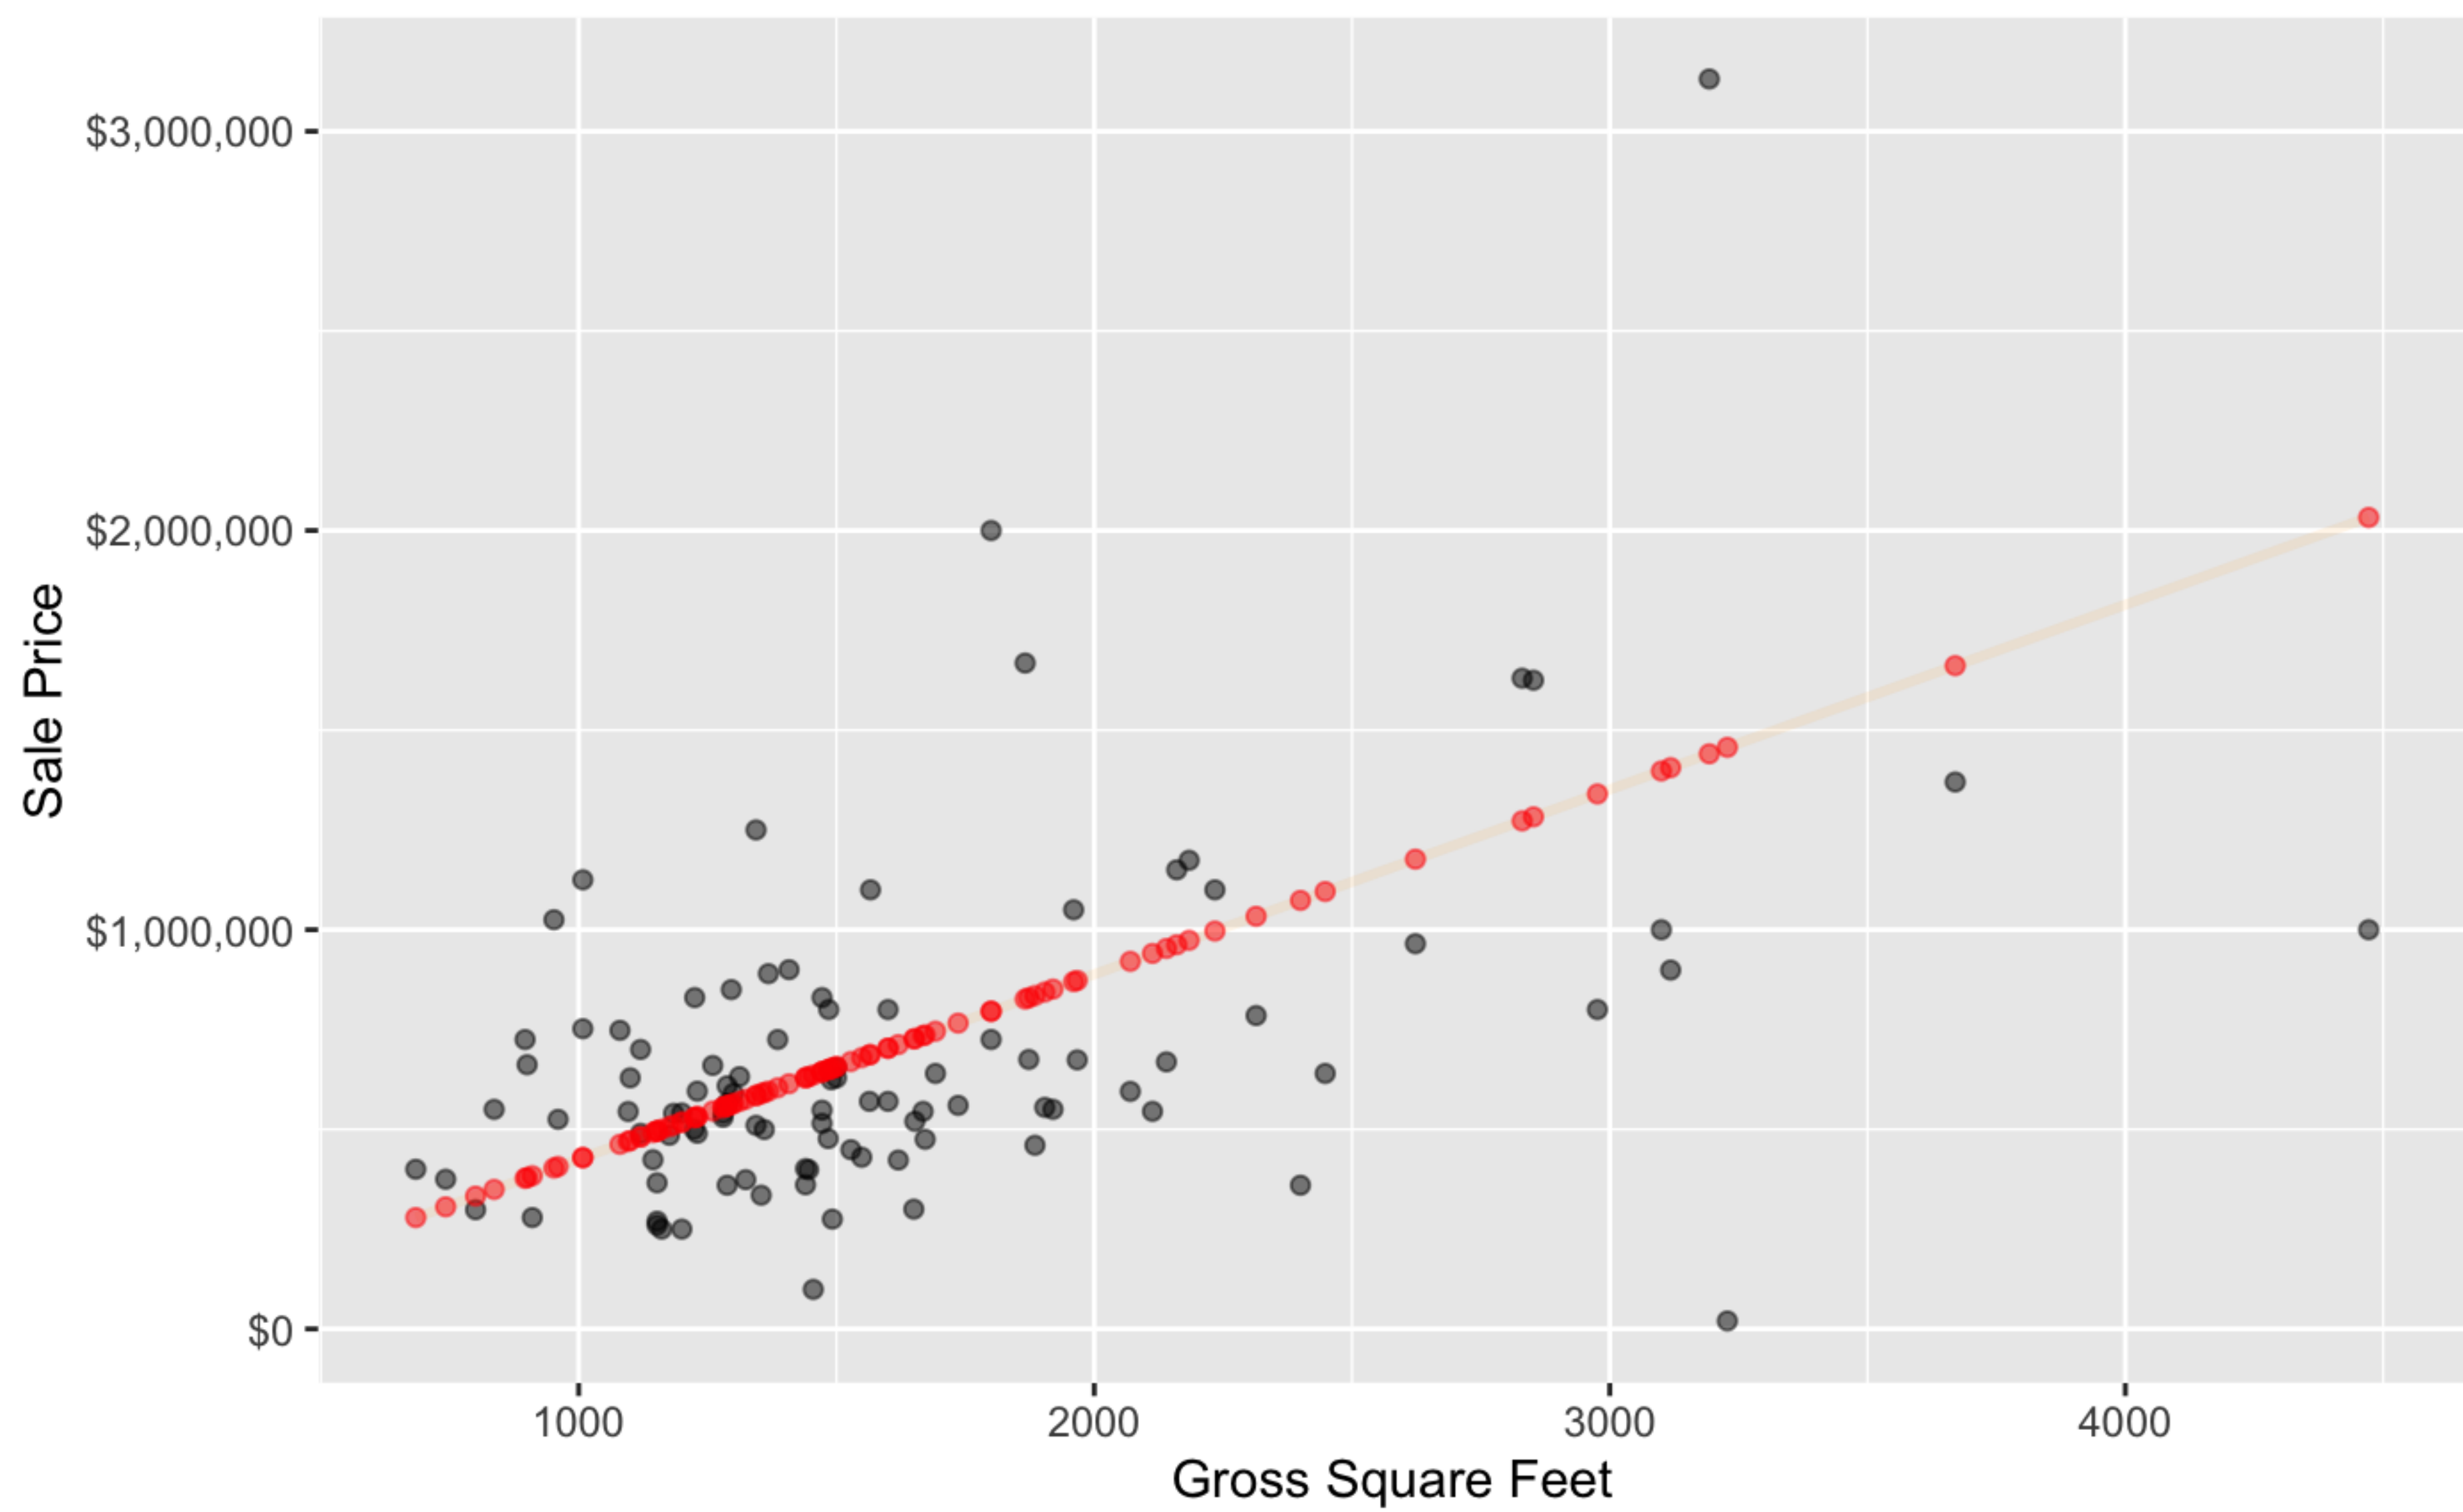
# More techniques


**Lasso and Ridge Regression**


**Bagging and Boosting**




















**We want to move the green line closer to the blue line but only with access to the green dots (training data)**

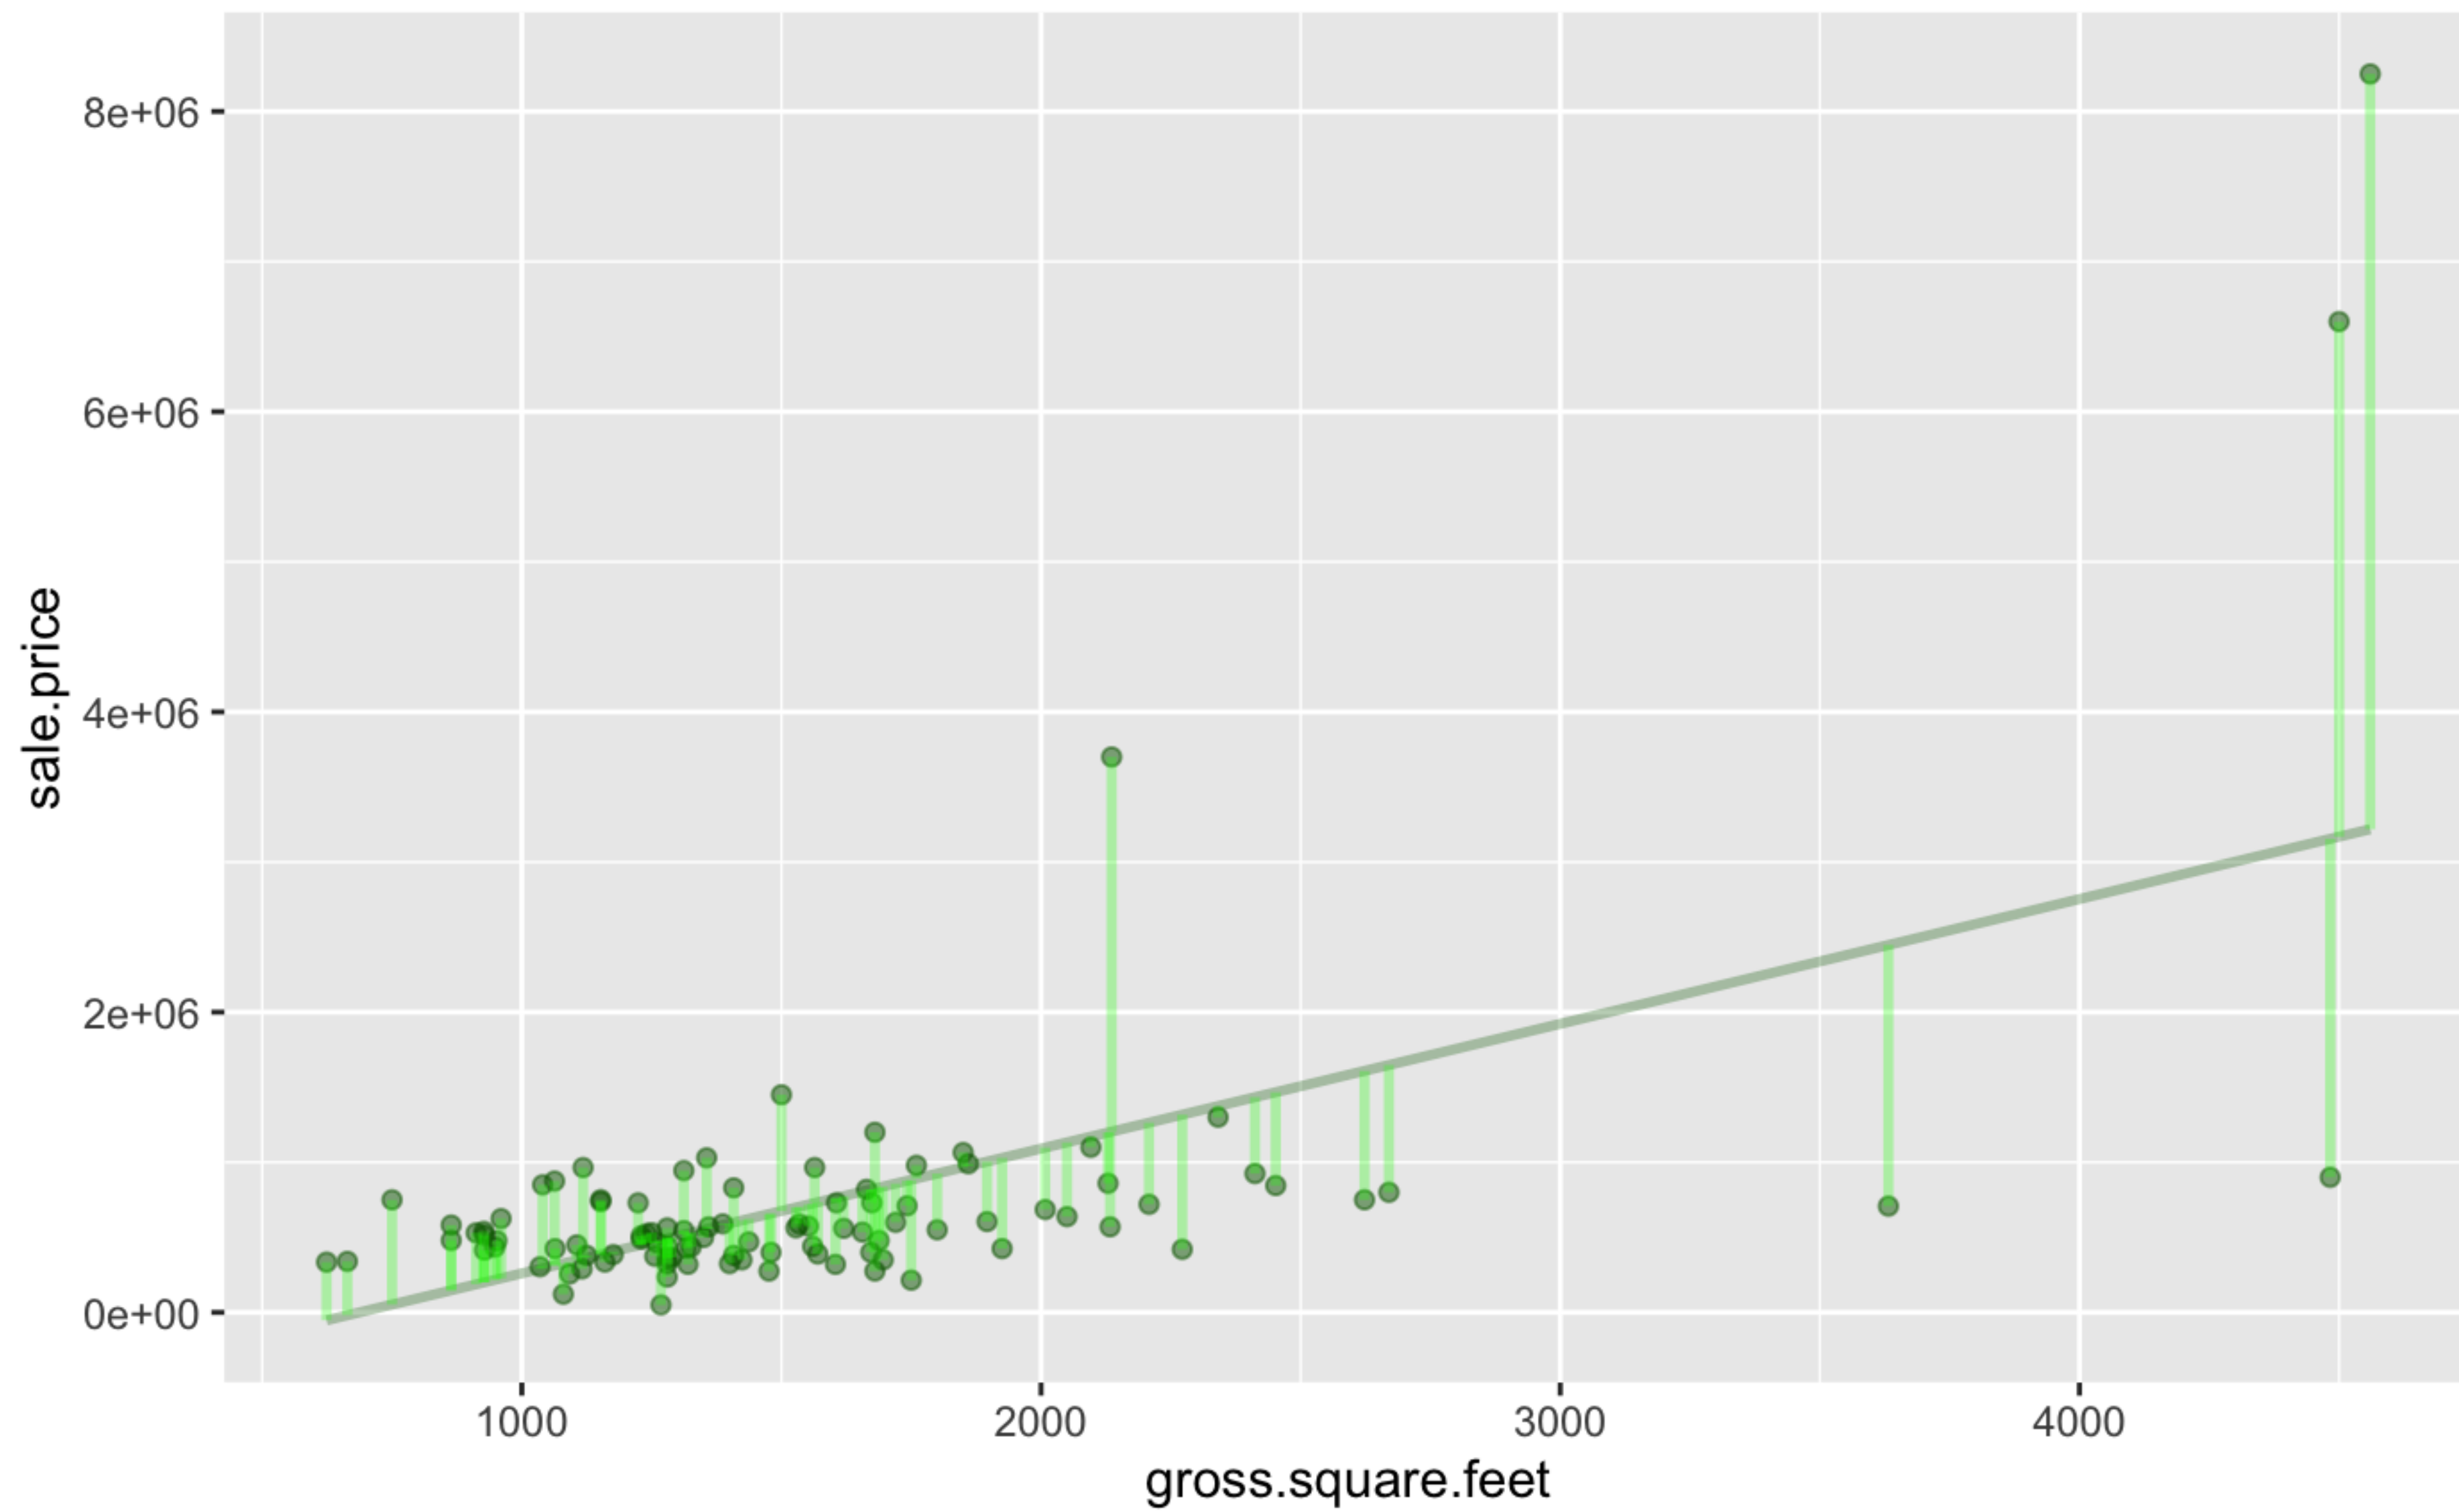
**That is, we want to reduce the  
overfitting to the training data**

**We want to REDUCE variance**


**We want to INCREASE bias**



$$sale . price = gross . square . feet * m + c$$



$$sale . price = gross . square . feet * m + c$$



$$sale . price = gross . square . feet * m + c$$

$$sale . price = gross . square . feet * m + c$$

- Linear Regression
  - Minimize “errors”
  - Minimize RMSE - equivalent to minimize “sum of the squared residuals”
- **Ridge Regression**
  - Minimize “sum of the squared residuals + **penalty**”
  - Minimize “sum of the squared residuals +  $\lambda * m^2$ ”





- **Ridge Regression**

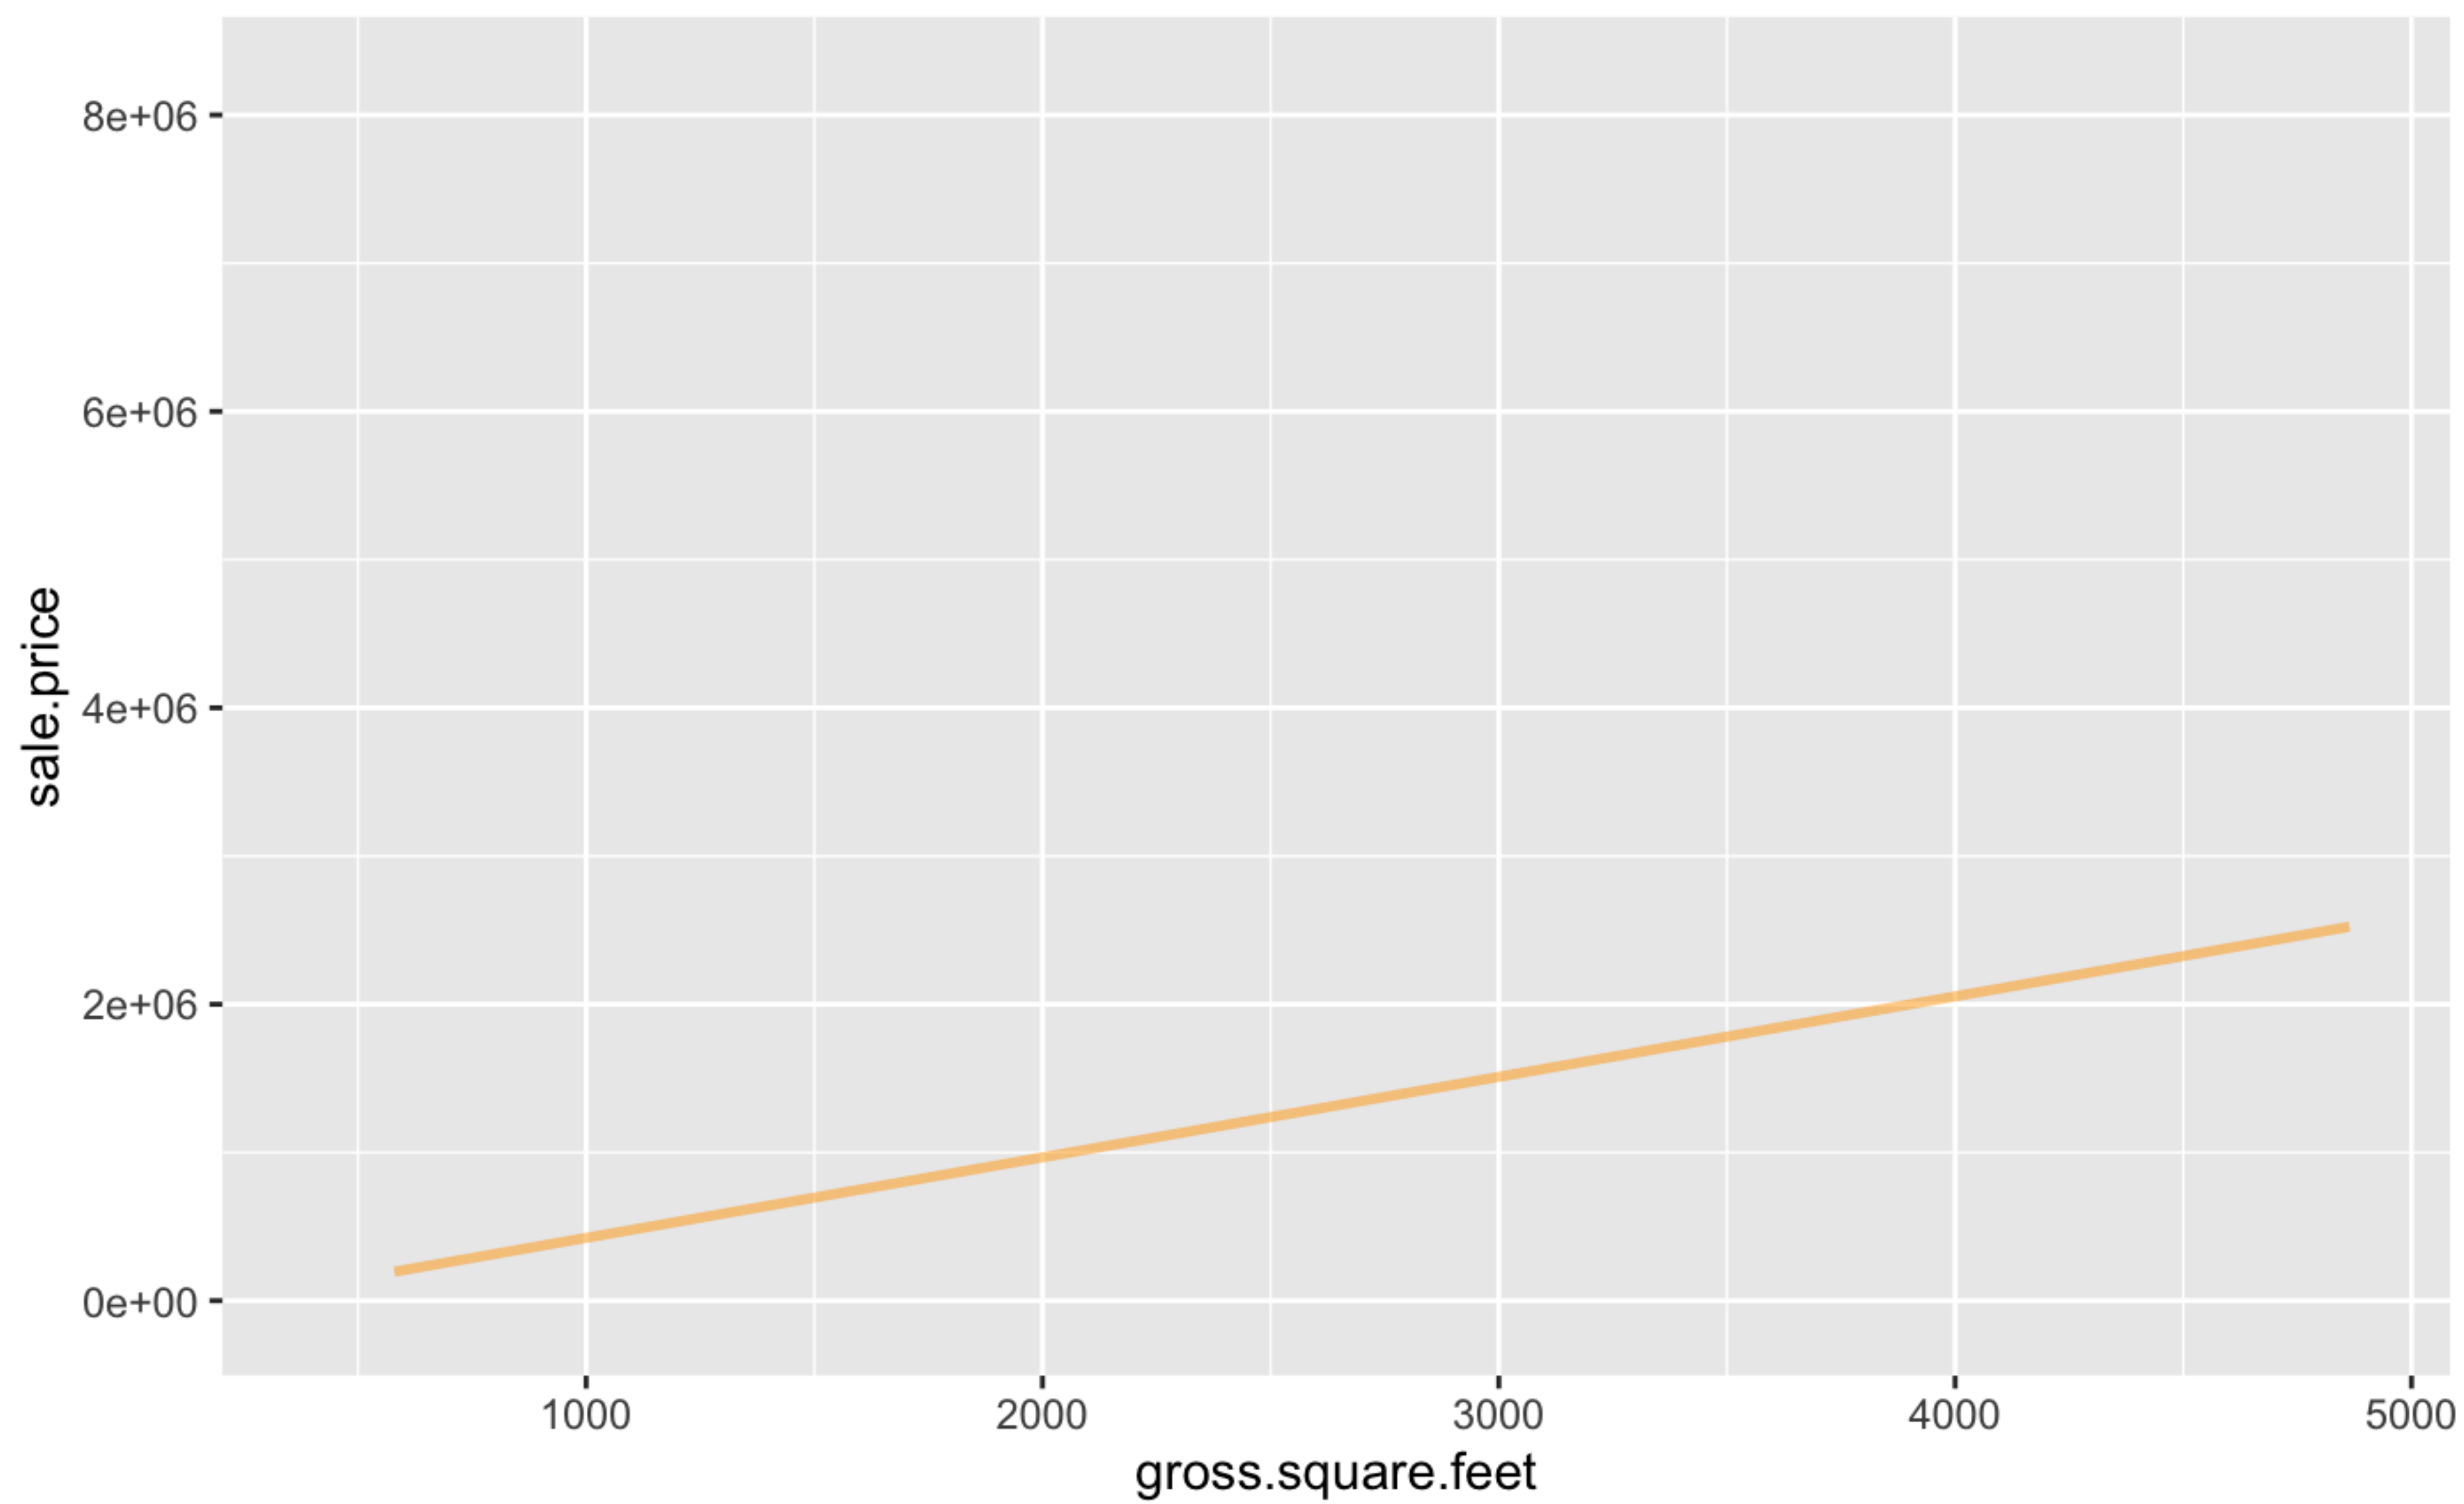
- Minimize “sum of the squared residuals + **penalty**”
- Minimize “sum of the squared residuals +  $\lambda * m^2$ ”
- Penalty gets smaller as  $m^2$  gets smaller, i.e. penalty is lower when slope shrinks towards zero - that is what we want!


# Why?


- Why do we want to shrink slope towards zero
- When slope is large - small change in  $\mathbf{x}$  creates large changes in  $\mathbf{y}$ 
  - High variance - sign of overfit
- Shrinking slope implies
  - small change in  $\mathbf{x}$  creates small(er) changes in  $\mathbf{y}$
  - Lower variance - cost of added bias











# What about the $\lambda$ ?

- $\lambda$  (lambda) can be any value from 0 -> infinity
- Larger  $\lambda$  means larger penalty -> slope closer to 0
- Smaller  $\lambda$  means lower penalty -> slope closer to normal linear regression
- If  $\lambda = 0$ , Ridge regression = linear regression !!!!!!!



How to decide  $\lambda$ ?  
Hyper parameter tuning!

```
592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
```

```
> ridge_spec
```

Linear Regression Model Specification (regression)

Main Arguments:

penalty = tune()

mixture = 0

Computational engine: glmnet

```
592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
```

```
> cv_folds
# 5-fold cross-validation
# A tibble: 5 × 2
  splits          id
  <list>         <chr>
1 <split [100/26]> Fold1
2 <split [101/25]> Fold2
3 <split [101/25]> Fold3
4 <split [101/25]> Fold4
5 <split [101/25]> Fold5
```

```

592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601

```

```

> lambda_grid%>%
+   filter(row_number() %% 5 == 1)
# A tibble: 20 × 1
  penalty
  <dbl>
1  0.00001
2  0.0000320
3  0.000102
4  0.000327
5  0.00105
6  0.00335
7  0.0107
8  0.0343
9  0.110
10 0.351
11 1.12
12 3.59
13 11.5
14 36.8
15 118.
16 376.
17 1205.
18 3854.
19 12328.
20 39442.

```

```
penalty(range = c(-10, 0), trans = transform_log10())
```

### Arguments

- range** A two-element vector holding the *defaults* for the smallest and largest possible values, respectively. If a transformation is specified, these values should be in the *transformed units*.
- trans** A trans object from the scales package, such as `scales::transform_log10()` or `scales::transform_reciprocal()`. If not provided, the default is used which matches the units used in range. If no transformation, NULL.

```
592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
```

```

592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
605 ridge_wf <- workflow() %>%
606   add_model(ridge_spec) %>%
607   add_recipe(ridge_recipe)
608

```

```
> ridge_wf
```

```
== Workflow ==
```

```
Preprocessor: Recipe
```

```
Model: linear_reg()
```

```
— Preprocessor —
```

```
0 Recipe Steps
```

```
— Model —
```

```
Linear Regression Model Specification (regression)
```

```
Main Arguments:
```

```
  penalty = tune()
```

```
  mixture = 0
```

```
Computational engine: glmnet
```



```

592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
605 ridge_wf <- workflow() %>%
606   add_model(ridge_spec) %>%
607   add_recipe(ridge_recipe)
608
609 # Tune the model
610 tune_results <- tune_grid(
611   ridge_wf,
612   resamples = cv_folds,
613   grid = lambda_grid
614 )
615
616 # Find best lambda
617 best_lambda <- select_best(tune_results, metric = "rmse")

```

```

> best_lambda
# A tibble: 1 × 2
  penalty .config
  <dbl> <chr>
1  100000 Preprocessor1_Model100

```

```
592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
605 ridge_wf <- workflow() %>%
606   add_model(ridge_spec) %>%
607   add_recipe(ridge_recipe)
608
609 # Tune the model
610 tune_results <- tune_grid(
611   ridge_wf,
612   resamples = cv_folds,
613   grid = lambda_grid
614 )
615
616 # Find best lambda
617 best_lambda <- select_best(tune_results, metric = "rmse")
618
619 # Finalize workflow with best lambda
620 final_ridge_workflow <- ridge_wf %>%
621   finalize_workflow(best_lambda)
622
623 # Fit final model
624 ridge_fit <- final_ridge_workflow %>%
625   fit(data = train)
626
627 # Look at results
628 ridge_fit %>%
629   tidy()
```



```

592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
605 ridge_wf <- workflow() %>%
606   add_model(ridge_spec) %>%
607   add_recipe(ridge_recipe)
608
609 # Tune the model
610 tune_results <- tune_grid(
611   ridge_wf,
612   resamples = cv_folds,
613   grid = lambda_grid
614 )
615
616 # Find best lambda
617 best_lambda <- select_best(tune_results, metric = "rmse")
618
619 # Finalize workflow with best lambda
620 final_ridge_workflow <- ridge_wf %>%
621   finalize_workflow(best_lambda)
622

```

```
> final_ridge_workflow
```

```
== Workflow ==
```

```
Preprocessor: Recipe
```

```
Model: linear_reg()
```

```
— Preprocessor —
```

```
0 Recipe Steps
```

```
— Model —
```

```
Linear Regression Model Specification (regression)
```


```
Main Arguments:
```

```
penalty = 1e+05
```

```
mixture = 0
```

```
Computational engine: glmnet
```

```
592 set.seed(345)
593 ridge_spec <- linear_reg(penalty = tune(), mixture = 0) %>%
594   set_engine("glmnet")
595
596 # Create cross validation folds
597 cv_folds <- vfold_cv(train, v = 5)
598
599 # Set up grid of lambda values to try
600 lambda_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 100)
601
602 train <- train %>% mutate(dummy = 1)
603 ridge_recipe <- recipe(sale.price ~ gross.square.feet + dummy, data = train)
604
605 ridge_wf <- workflow() %>%
606   add_model(ridge_spec) %>%
607   add_recipe(ridge_recipe)
608
609 # Tune the model
610 tune_results <- tune_grid(
611   ridge_wf,
612   resamples = cv_folds,
613   grid = lambda_grid
614 )
615
616 # Find best lambda
617 best_lambda <- select_best(tune_results, metric = "rmse")
618
619 # Finalize workflow with best lambda
620 final_ridge_workflow <- ridge_wf %>%
621   finalize_workflow(best_lambda)
622
623 # Fit final model
624 ridge_fit <- final_ridge_workflow %>%
625   fit(data = train)
626
627 # Look at results
628 ridge_fit %>%
629   tidy()
```



```
652 rmse_ridge <- rmse(lm_ridge_test_predictions, truth = sale.price, estimate = ridge_pred)
653 lmse_ridge <- rmse(lm_ridge_test_predictions, truth = sale.price, estimate = lm_pred)
654
655 print(paste("Ridge RMSE:", round(rmse_ridge$.estimate)))
656 print(paste("LM RMSE:", round(lmse_ridge$.estimate)))
657 ```
```

[1] "Ridge RMSE: 490396"

[1] "LM RMSE: 507865"

# Notes

- With multiple predictors the error is  $\lambda(m_1^2 + m_2^2 + \dots + m_n^2)$
- Works on Logistic regression as well
- Mixture = 0 -> Ridge regression
- Mixture = 1 -> Lasso regression

# Lasso Regression

- Lasso works similarly but uses a different penalty
- $\lambda * (|m_1| + |m_2| + \dots + |m_n|)$

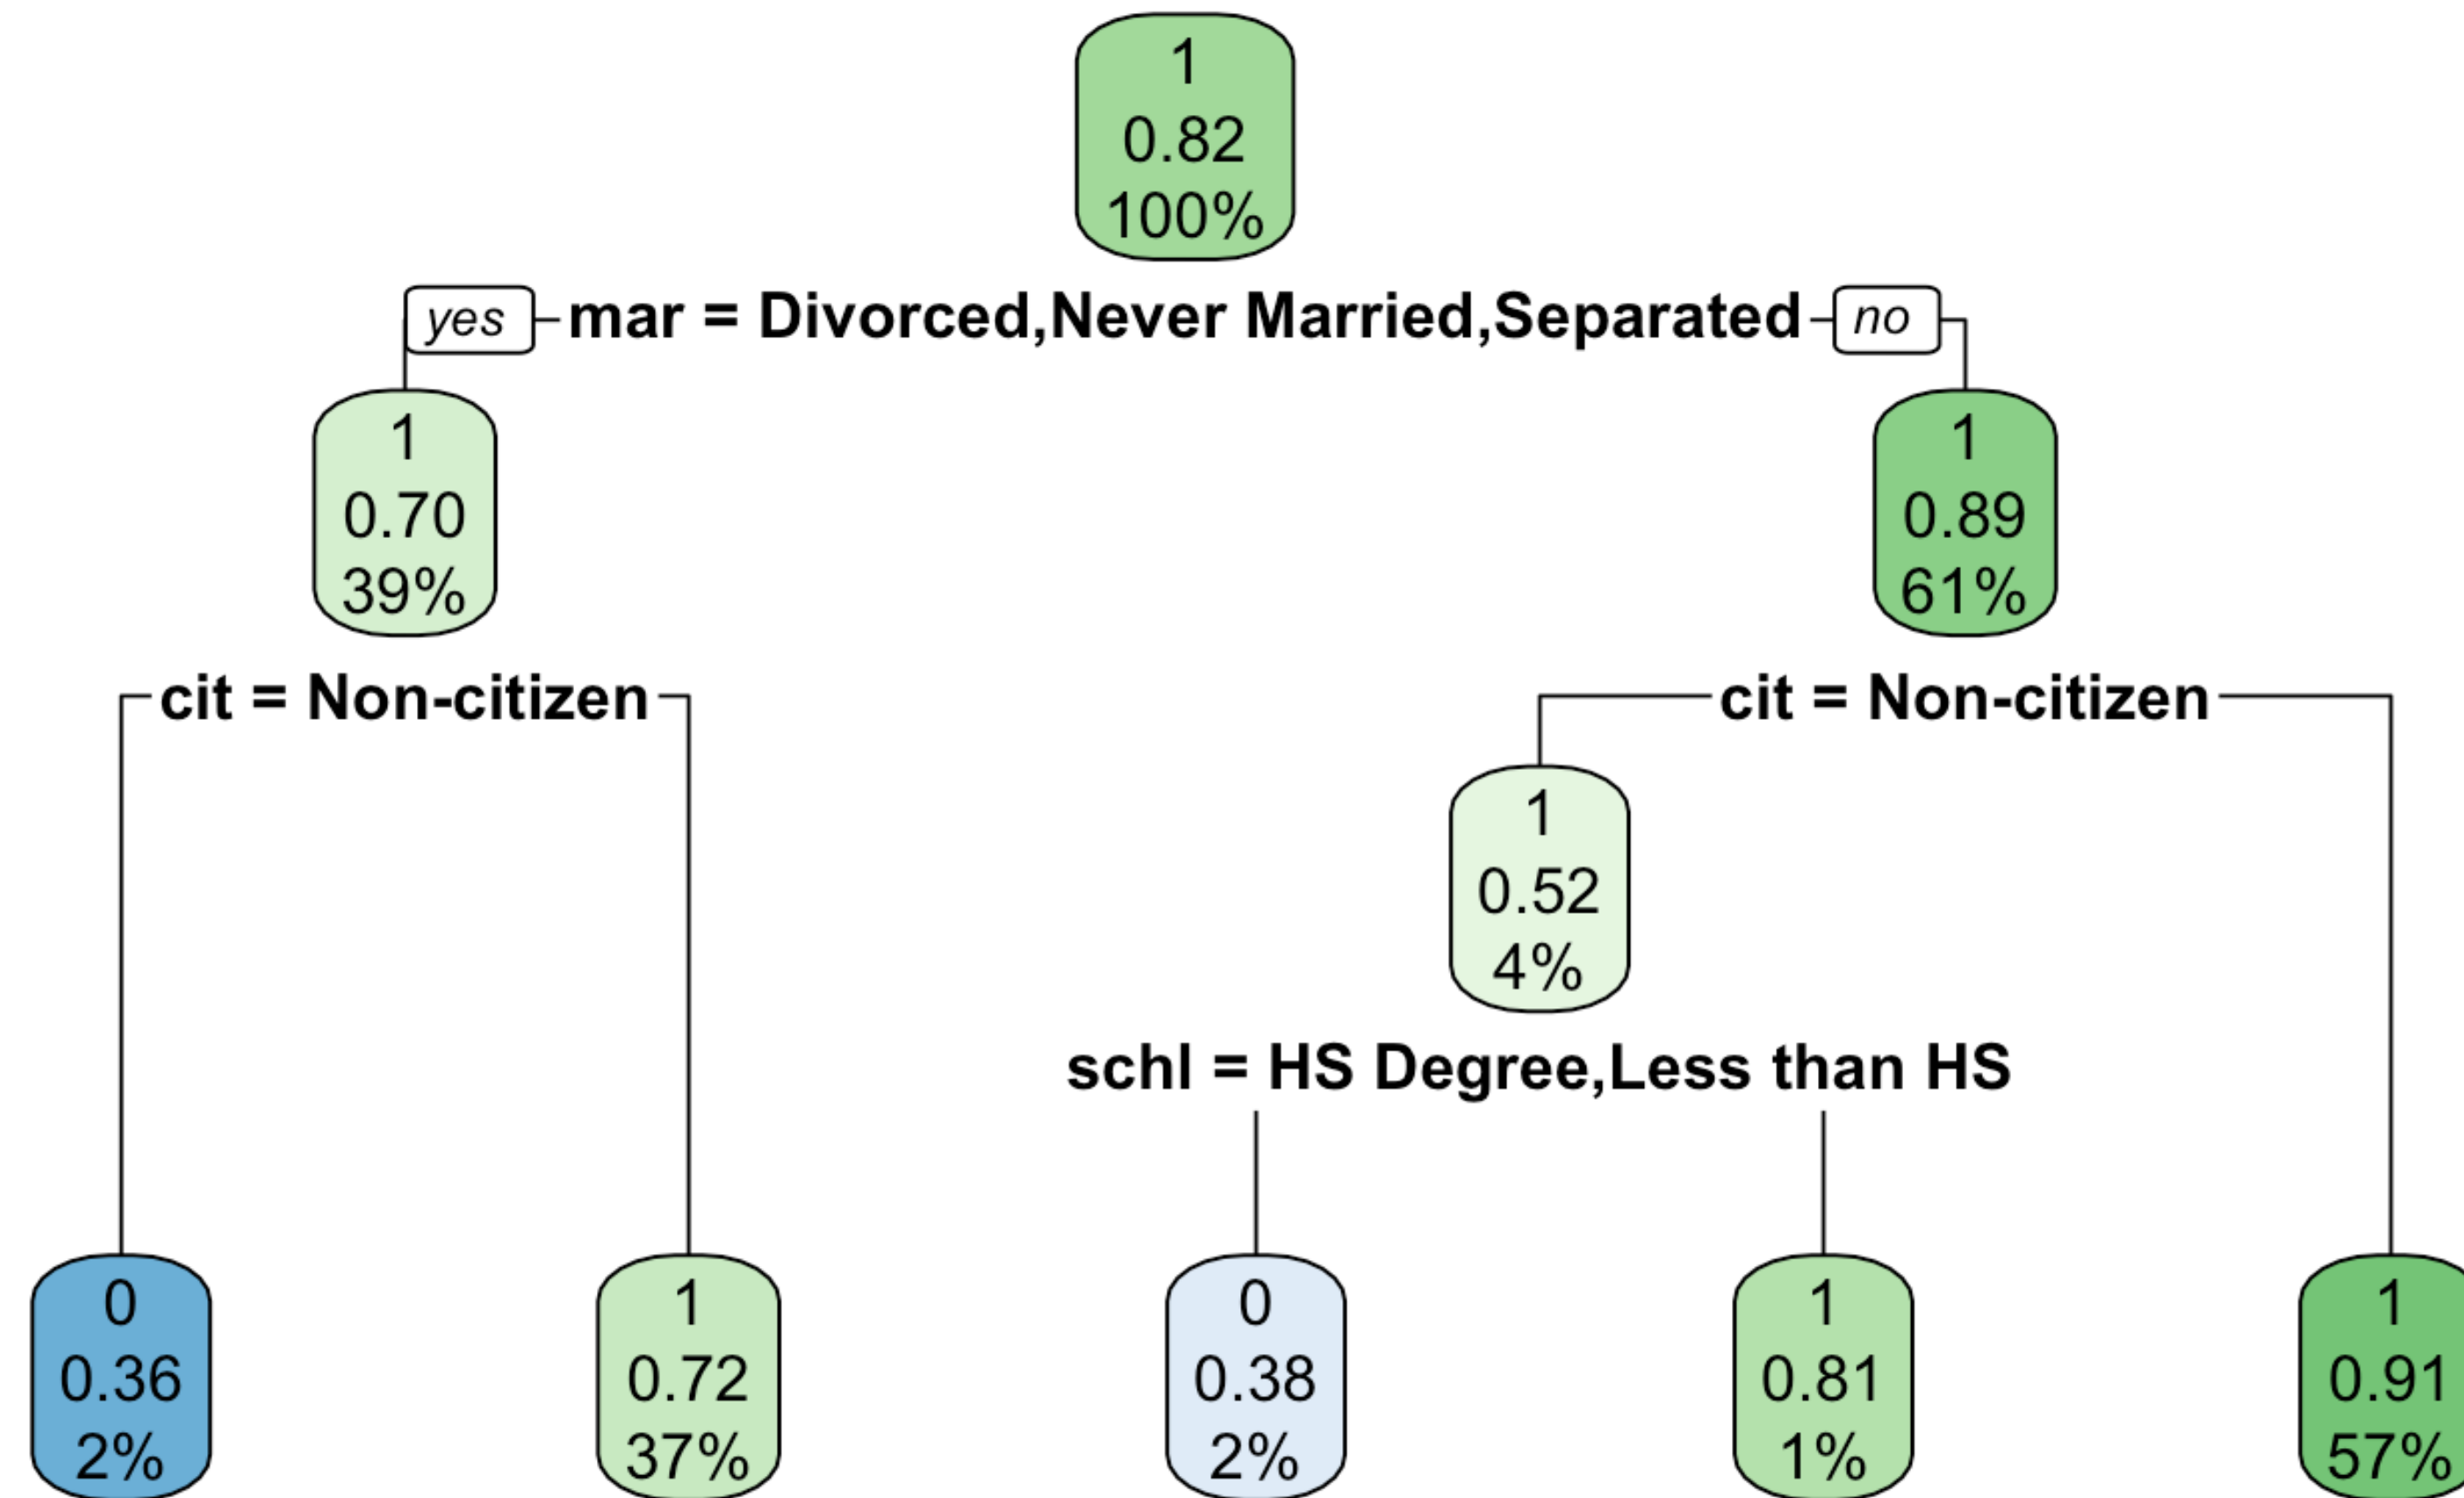
# Lasso vs Ridge

- Ridge -
  - “L2 Norm” - Euclidean distance
  - Moves closer to 0
  - Use when all variables are relevant but need shrinkage to avoid overfitting.
- Lasso
  - “L1 Norm” - Manhattan distance
  - Can set to exactly 0 - useful for variable selection
  - Use when you suspect some features are irrelevant for feature selection

# Decision Tree



R Console





# Bagging (Bootstrap AGGregating)

- Uses all available features when building each decision tree. It doesn't perform feature randomization at each split.
- Since all features are considered at each split, the trees in a bagged ensemble tend to be more correlated with each other.



- Difference from random\_forest —
  - Randomly select a subset of features at each node split in each tree
- Similarity —
  - Both use random subsets of data for each tree that is trained

# Boosting

- Boosting is also an ensemble technique that builds trees sequentially.
  - We make a simple tree, see where it is weakest, and make that part better
- Process:
  - Train first tree on original data
  - Give higher weight to misclassified samples
  - Train next tree focusing on harder examples
  - Repeat, creating a chain of complementary trees

